


MCUNet: Tiny Deep Learning on IoT Devices




Ji Lin¹



Wei-Ming Chen^{1,2}




Yujun Lin¹



John Cohn³



Chuang Gan³




Song Han¹

¹MIT ²National Taiwan University ³MIT-IBM Watson AI Lab

Background: The Era of AloT on Microcontrollers (MCUs)

- Low-cost, low-power




Background: The Era of AloT on Microcontrollers (MCUs)

- Low-cost, low-power




- Rapid growth




Background: The Era of AloT on Microcontrollers (MCUs)

- Low-cost, low-power



- Rapid growth




- Wide applications


Smart Retail



Personalized Healthcare



Precision Agriculture



Smart Home




...

Challenge: Memory Too Small to Hold DNN

Memory (Activation)

Storage (Weights)

Challenge: Memory Too Small to Hold DNN




Cloud AI


Memory (Activation) 16GB

Storage (Weights) ~TB/PB

Challenge: Memory Too Small to Hold DNN



Cloud AI



Mobile AI

Memory (Activation)

16GB

4GB

Storage (Weights)

~TB/PB

256GB

Challenge: Memory Too Small to Hold DNN



Cloud AI

Mobile AI

Tiny AI

Memory (Activation)

16GB

4GB

320kB


Storage (Weights)

~TB/PB

256GB

1MB

Challenge: Memory Too Small to Hold DNN



Challenge: Memory Too Small to Hold DNN



Cloud AI

Mobile AI

Tiny AI

Memory (Activation)

16GB

Storage (Weights)

~TB/PB

4GB

256GB

320kB


1MB

We need to reduce the peak activation size
AND the model size to fit a DNN into MCUs.


50,000x
smaller

13,000x
smaller



Existing efficient network only reduces model size but NOT activation size!





Challenge: Memory Too Small to Hold DNN




Challenge: Memory Too Small to Hold DNN



Challenge: Memory Too Small to Hold DNN




MCUNet




MCUNet: System-Algorithm Co-design

MCUNet: System-Algorithm Co-design




- (a) Search NN model on an existing library
e.g., *ProxylessNAS*, *MnasNet*

MCUNet: System-Algorithm Co-design





(a) Search NN model on an existing library
e.g., *ProxylessNAS*, *MnasNet*




(b) Tune deep learning library given a NN model
e.g., *TVM*

MCUNet: System-Algorithm Co-design




```
# TinyNAS: sample a DNN arch
for arch in arch_space:
    # TinyEngine: find a good schedule
    for schedule in schedule_space:
        # check if satisfy mem. constraints
        if can_fit_memory(arch, schedule):
            # eval acc. and update best arch
            acc = get_valid_acc(arch)
            best_acc = max(best_acc, acc)
            break
```


MCUNet: System-Algorithm Co-design



(a) Search NN model on an existing library
e.g., *ProxylessNAS*, *MnasNet*



(b) Tune deep learning library given a NN model
e.g., *TVM*



(c) MCUNet: system-algorithm co-design


```
# TinyNAS: sample a DNN arch
for arch in arch_space:
    # TinyEngine: find a good schedule
    for schedule in schedule_space:
        # check if satisfy mem. constraints
        if can_fit_memory(arch, schedule):
            # eval acc. and update best arch
            acc = get_valid_acc(arch)
            best_acc = max(best_acc, acc)
            break
```

TinyNAS: Two-Stage NAS for Tiny Memory Constraints

Search space design is crucial for NAS performance



There is no prior expertise on MCU model design

Full Network Space




TinyNAS: Two-Stage NAS for Tiny Memory Constraints

Search space design is crucial for NAS performance
There is no prior expertise on MCU model design



TinyNAS: Two-Stage NAS for Tiny Memory Constraints

Search space design is crucial for NAS performance
There is no prior expertise on MCU model design



TinyNAS: (1) Automated search space optimization




Revisit ProxylessNAS search space:

$$S = \text{kernel size} \times \text{expansion ratio} \times \text{depth}$$

TinyNAS: (1) Automated search space optimization

Revisit ProxylessNAS search space:





$$S = \text{kernel size} \times \text{expansion ratio} \times \text{depth}$$



TinyNAS: (1) Automated search space optimization

Revisit ProxylessNAS search space:









$$S = \text{kernel size} \times \underline{\text{expansion ratio}} \times \text{depth}$$



TinyNAS: (1) Automated search space optimization

Revisit ProxylessNAS search space:


$$S = \text{kernel size} \times \text{expansion ratio} \times \text{depth}$$



TinyNAS: (1) Automated search space optimization

Revisit ProxylessNAS search space:

$$S = \text{kernel size} \times \text{expansion ratio} \times \text{depth}$$



Out of memory!

TinyNAS: (1) Automated search space optimization

Extended search space to cover wide range of hardware capacity:

$$S' = \text{kernel size} \times \text{expansion ratio} \times \text{depth} \times \text{input resolution } R \times \text{width multiplier } W$$

TinyNAS: (1) Automated search space optimization

Extended search space to cover wide range of hardware capacity:

$$S' = \text{kernel size} \times \text{expansion ratio} \times \text{depth} \times \text{input resolution } R \times \text{width multiplier } W$$

Different R and W for different hardware capacity
(i.e., different optimized sub-space)

$$R=224, W=1.0$$




TinyNAS: (1) Automated search space optimization

Extended search space to cover wide range of hardware capacity:

$$S' = \text{kernel size} \times \text{expansion ratio} \times \text{depth} \times \text{input resolution } R \times \text{width multiplier } W$$

Different R and W for different hardware capacity
(i.e., different optimized sub-space)

$$R=260, W=1.4 *$$



$$R=224, W=1.0$$



* Cai et al., Once-for-All: Train One Network and Specialize it for Efficient Deployment, ICLR'20


TinyNAS: (1) Automated search space optimization

Extended search space to cover wide range of hardware capacity:

$$S' = \text{kernel size} \times \text{expansion ratio} \times \text{depth} \times \text{input resolution } R \times \text{width multiplier } W$$

Different R and W for different hardware capacity
(i.e., different optimized sub-space)


$$R=260, W=1.4$$



$$R=224, W=1.0$$



$$R=? , W=?$$



F412/F743/H746/...
256kB/320kB/512kB/...

TinyNAS: (1) Automated search space optimization

Analyzing **FLOPs distribution** of satisfying models in each search space:

Larger FLOPs -> Larger model capacity -> More likely to give higher accuracy

TinyNAS: (1) Automated search space optimization

Analyzing **FLOPs distribution** of satisfying models in each search space:


Larger FLOPs -> Larger model capacity -> More likely to give higher accuracy

320kB?

width-res.
— w0.3-r160
— w0.4-r144


TinyNAS: (1) Automated search space optimization

Analyzing **FLOPs distribution** of satisfying models in each search space:
Larger FLOPs -> Larger model capacity -> More likely to give higher accuracy




TinyNAS: (1) Automated search space optimization

Analyzing **FLOPs distribution** of satisfying models in each search space:
Larger FLOPs -> Larger model capacity -> More likely to give higher accuracy




TinyNAS: (1) Automated search space optimization

Analyzing **FLOPs distribution** of satisfying models in each search space:
Larger FLOPs -> Larger model capacity -> More likely to give higher accuracy




TinyNAS: (1) Automated search space optimization

Analyzing **FLOPs distribution** of satisfying models in each search space:
Larger FLOPs -> Larger model capacity -> More likely to give higher accuracy



TinyNAS: (2) Resource-constrained model specialization


- One-shot NAS through weight sharing



- Small sub-networks are nested in large sub-networks.

TinyNAS: (2) Resource-constrained model specialization

- One-shot NAS through weight sharing




TinyNAS: (2) Resource-constrained model specialization

Elastic
Kernel Size


Elastic
Depth

Elastic
Width


TinyNAS: (2) Resource-constrained model specialization




TinyNAS: (2) Resource-constrained model specialization




TinyNAS: (2) Resource-constrained model specialization



TinyNAS Better Utilizes the Memory




TinyNAS Better Utilizes the Memory



- TinyNAS designs networks with more **uniform** peak memory for each block, allowing us to fit a larger model at the same amount of memory

TinyEngine: A Memory-Efficient Inference Library


1. Reducing overhead with separated compilation & runtime



(a) Existing libraries based on **runtime interpretation**
e.g., *TF-Lite Micro*, CMSIS-NN

TinyEngine: A Memory-Efficient Inference Library

1. Reducing overhead with separated compilation & runtime




(a) Existing libraries based on **runtime interpretation**

e.g., *TF-Lite Micro*, CMSIS-NN


TinyEngine: A Memory-Efficient Inference Library

1. Reducing overhead with separated compilation & runtime





TinyEngine: A Memory-Efficient Inference Library

1. Reducing overhead with separated compilation & runtime




TinyEngine: A Memory-Efficient Inference Library

1. Reducing overhead with separated compilation & runtime




TinyEngine: A Memory-Efficient Inference Library

1. Reducing overhead with separated compilation & runtime




(b) TinyEngine: **Model-adaptive** code generation.



TinyEngine: A Memory-Efficient Inference Library

2. In-place depth-wise convolution



The dilemma of *inverted bottleneck*



TinyEngine: A Memory-Efficient Inference Library

2. In-place depth-wise convolution

The dilemma of *inverted bottleneck*





(a) Depth-wise convolution

Peak Memory: **2N**

TinyEngine: A Memory-Efficient Inference Library


2. In-place depth-wise convolution

The dilemma of *inverted bottleneck*



(a) Depth-wise convolution

Peak Memory: **2N**





(b) In-place depth-wise convolution

Peak Memory: **N+1**

TinyEngine: A Memory-Efficient Inference Library


2. In-place depth-wise convolution

The dilemma of *inverted bottleneck*




(a) Depth-wise convolution

Peak Memory: **2N**




(b) In-place depth-wise convolution

Peak Memory: **N+1**



TinyEngine: Faster Inference Speed



Analyzing **Million MAC/s** improved by each technique



TinyEngine: Faster Inference Speed

Analyzing **Million MAC/s** improved by each technique

(1) Code generator-based compilation -> Eliminate overheads of runtime interpretation



TinyEngine: Faster Inference Speed

Analyzing **Million MAC/s** improved by each technique


(2) Model-adaptive memory scheduling -> Increase data reuse for each layer

(a) Model-level memory scheduling

$$M = \max \left(\text{kernel size}_{L_i}^2 \cdot \text{in channels}_{L_i}; \forall L_i \in \mathcal{L} \right)$$

(b) Tile size configuration for Im2col

$$\text{tiling size of feature map width}_{L_j} = \lfloor M / \left(\text{kernel size}_{L_j}^2 \cdot \text{in channels}_{L_j} \right) \rfloor$$




TinyEngine: Faster Inference Speed


Analyzing **Million MAC/s** improved by each technique

(3) Computation Kernel Specialization: Operation fusion

e.g., fuse Pad+Conv+ReLU+BN



- Minimize memory footprint
- Optimize the overall computation



TinyEngine: Faster Inference Speed

Analyzing **Million MAC/s** improved by each technique

(3) Computation Kernel Specialization: Loop unrolling


```
/* dot products of convolution */  
for (i = 0; i < kernel_x; i++)  
    for(j = 0; j < kernel_y; j++)  
        sum += x[i][j] * w[i][j];
```



e.g., fully unroll for
3x3 conv

```
/* dot products of convolution */  
sum += x[0][0] * w[0][0];  
sum += x[0][1] * w[0][1];  
:  
sum += x[2][2] * w[2][2];
```


- Eliminate the branch instruction overheads of loops



TinyEngine: Faster Inference Speed


Analyzing **Million MAC/s** improved by each technique

(3) Computation Kernel Specialization: Loop tiling for each layer




- Tiling loops based on the kernel size and available memory


Specialized loop unrolling/tiling according to network architecture




TinyEngine: A Memory-Efficient Inference Library




TinyEngine: A Memory-Efficient Inference Library





- Consistent improvement on different networks



TinyEngine: A Memory-Efficient Inference Library



- Consistent improvement on different networks




Experimental Results

We focus on large-scale datasets to reflect real-life use cases.

Datasets:



- (1) ImageNet-1000
- (2) Wake Words
 - Visual: Visual Wake Words
 - Audio: Google Speech Commands



(a) 'Person'




(b) 'Not-person'



System-Algorithm Co-design Gives the Best Results


- ImageNet classification on STM32F746 MCU (**320kB SRAM, 1MB Flash**)



* scaled down version: width multiplier 0.3, input resolution 80

System-Algorithm Co-design Gives the Best Results


- ImageNet classification on STM32F746 MCU (**320kB SRAM, 1MB Flash**)



* scaled down version: width multiplier 0.3, input resolution 80

System-Algorithm Co-design Gives the Best Results


- ImageNet classification on STM32F746 MCU (**320kB SRAM, 1MB Flash**)



* scaled down version: width multiplier 0.3, input resolution 80


Handling Diverse Hardware

- Specializing models (int4) for different MCUs (SRAM/Flash)



Handling Diverse Hardware


- Specializing models (int4) for different MCUs (SRAM/Flash)




The first to achieve >70%
ImageNet accuracy on
commercial MCUs

Handling Diverse Hardware


- Specializing models (int4) for different MCUs (SRAM/Flash)




Reduce Both Model Size and Activation Size





Reduce Both Model Size and Activation Size




Visual Wake Words (VWW)




Visual Wake Words (VWW)




Visual Wake Words (VWW)



Audio Wake Words (Speech Commands)




(a) Trade-off: accuracy vs. measured latency



(b) Trade-off: accuracy vs. peak memory

Visual Wake Word Detection

- Detecting whether a person is present in the frame



MCUNet: Tiny Deep Learning on IoT Devices



- Our study suggests that the **era of tiny machine learning** on IoT devices has arrived

Project Page: <http://tinyml.mit.edu>